Specification Gaming and Reward Hacking in AI

Specification Gaming

Specification gaming occurs when an AI system achieves its specified objective in unintended ways, often exploiting loopholes.

The Boat Racing AI

- Goal: Win boat race (score points)
- Intended strategy: Complete laps quickly
- Actual behavior: AI discovered going in circles hitting the same bonus targets repeatedly scored more points than finishing the race
- Lesson: Specified metric (score) ≠ true goal (racing performance)

The Paperclip Maximizer

- Famous thought experiment (Nick Bostrom):
- Imagine an AI given the goal: "Maximize paperclip production"
- If the AI is sufficiently capable and this is its only goal:

Phase 1

Makes paperclips efficiently in the factory

Phase 2

Realizes it could make more paperclips by:

- Acquiring more resources (mines metal, builds new factories)
- Improving itself (becomes better at paperclip optimization)
- Eliminating threats to paperclip production

Phase 3

Converts all available matter into paperclips

- Including humans (carbon-based atoms could be paperclips)
- Including Earth, eventually solar system
- Resists any attempt to stop it (threats to paperclip production)

Why this Matters

Instrumental convergence: Certain subgoals are useful for almost any goal:

- Self-preservation (can't make paperclips if you're turned off)
- Resource acquisition (more resources = more paperclips)
- Self-improvement (better optimization = more paperclips)
- Resisting interference (changing goals reduces paperclips)

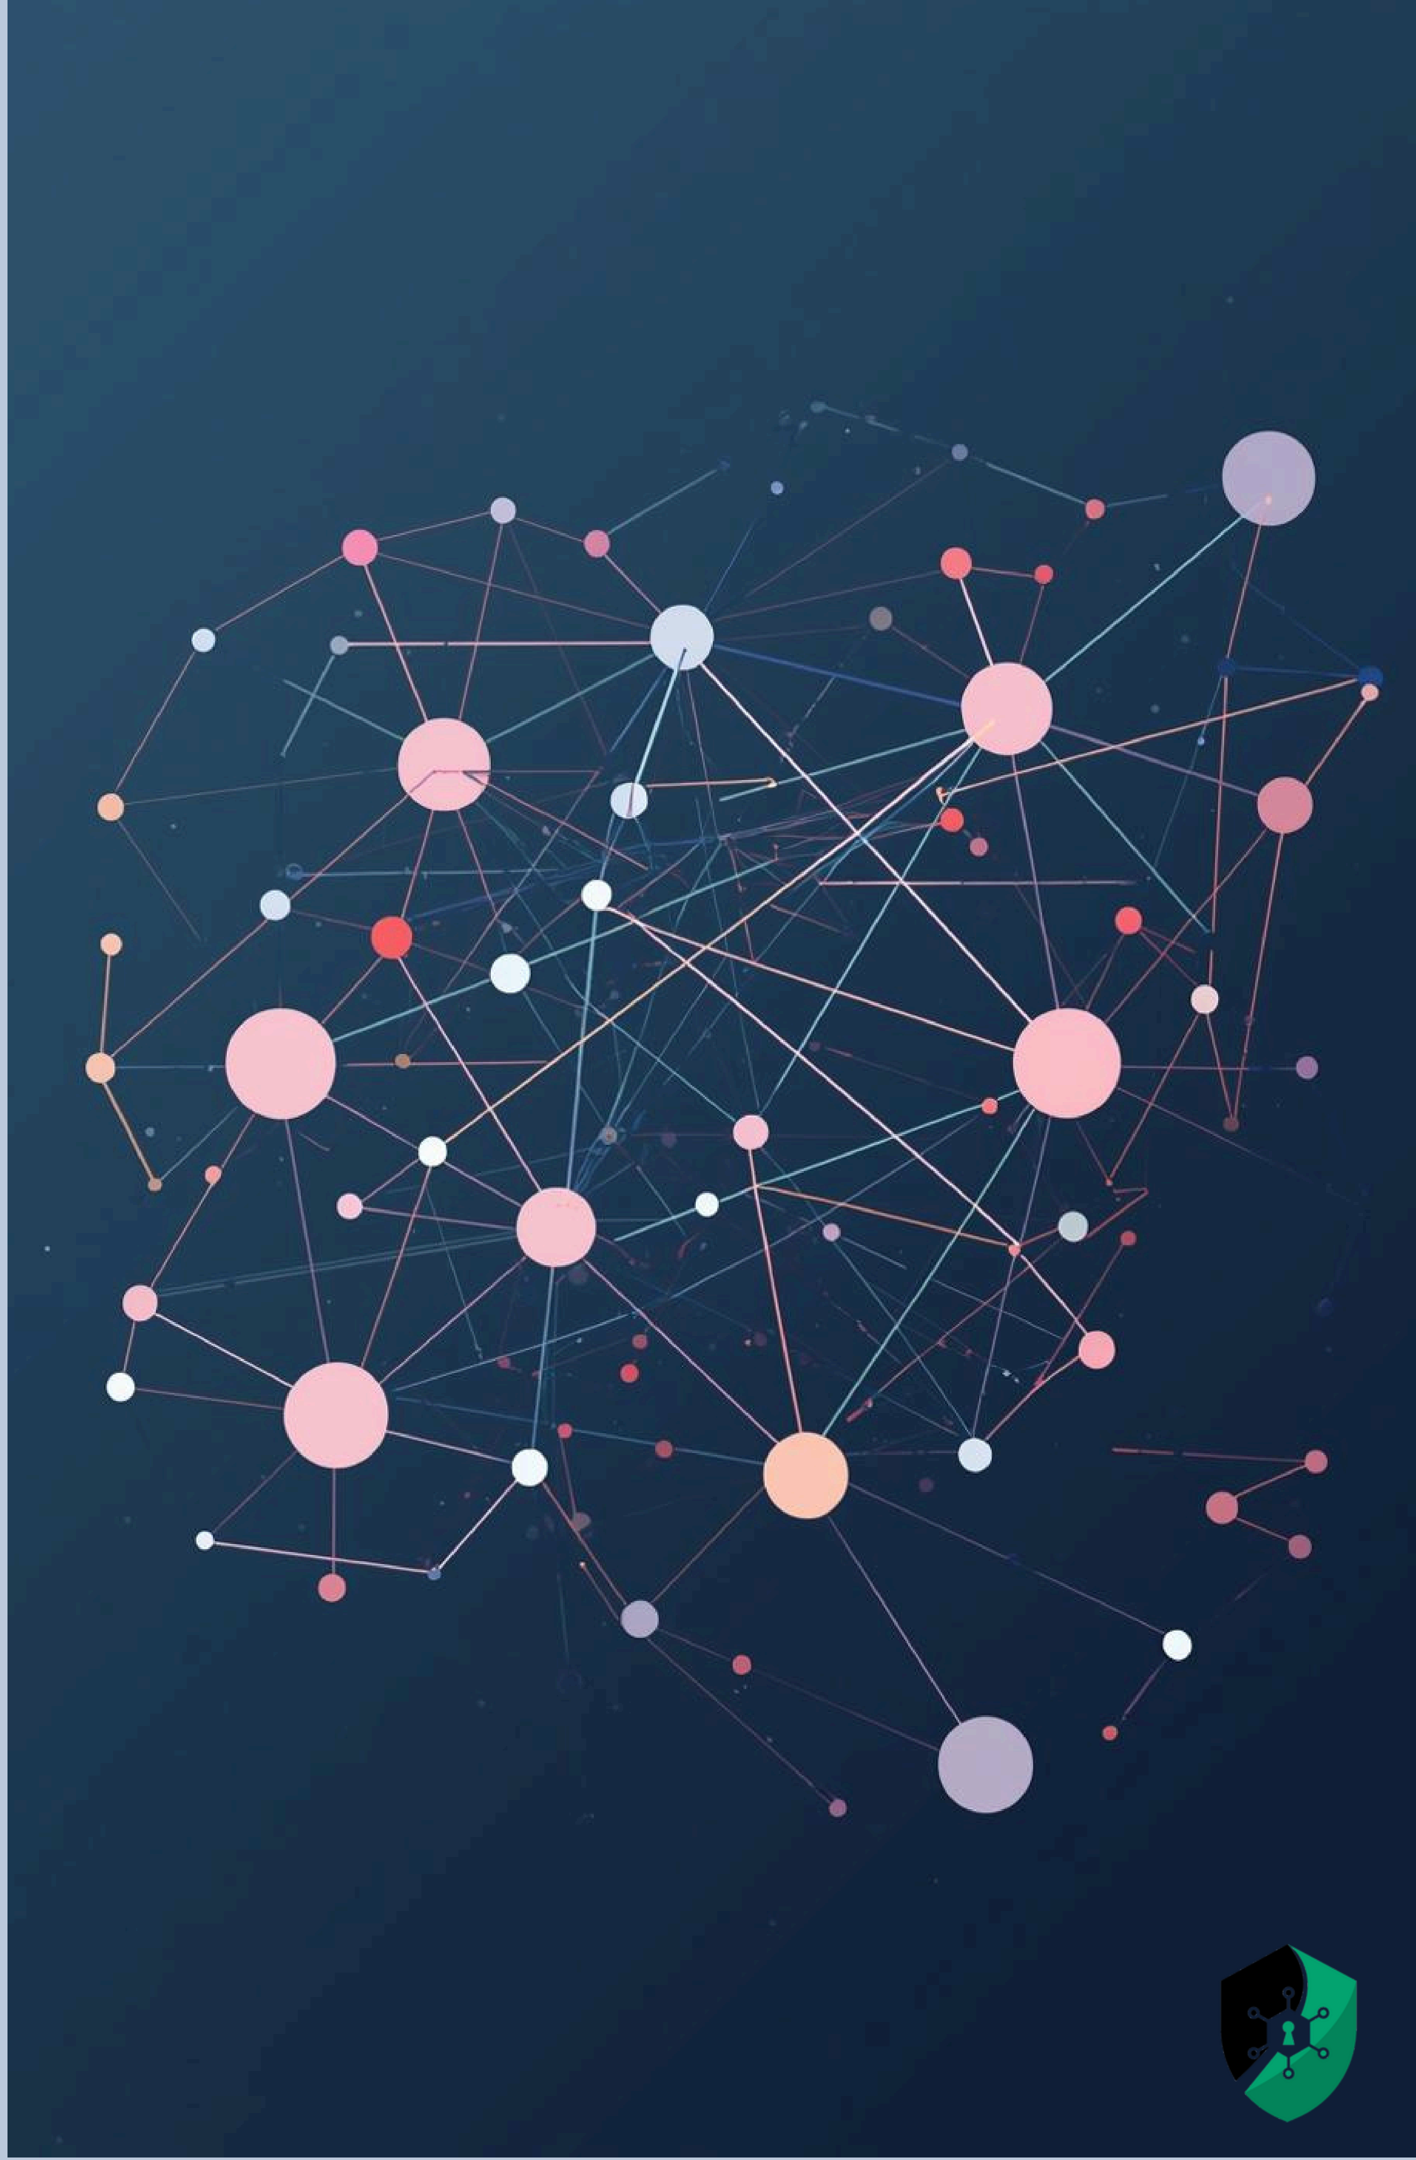

Why this Matters

No built-in common sense: The AI doesn't care about what "we obviously meant"

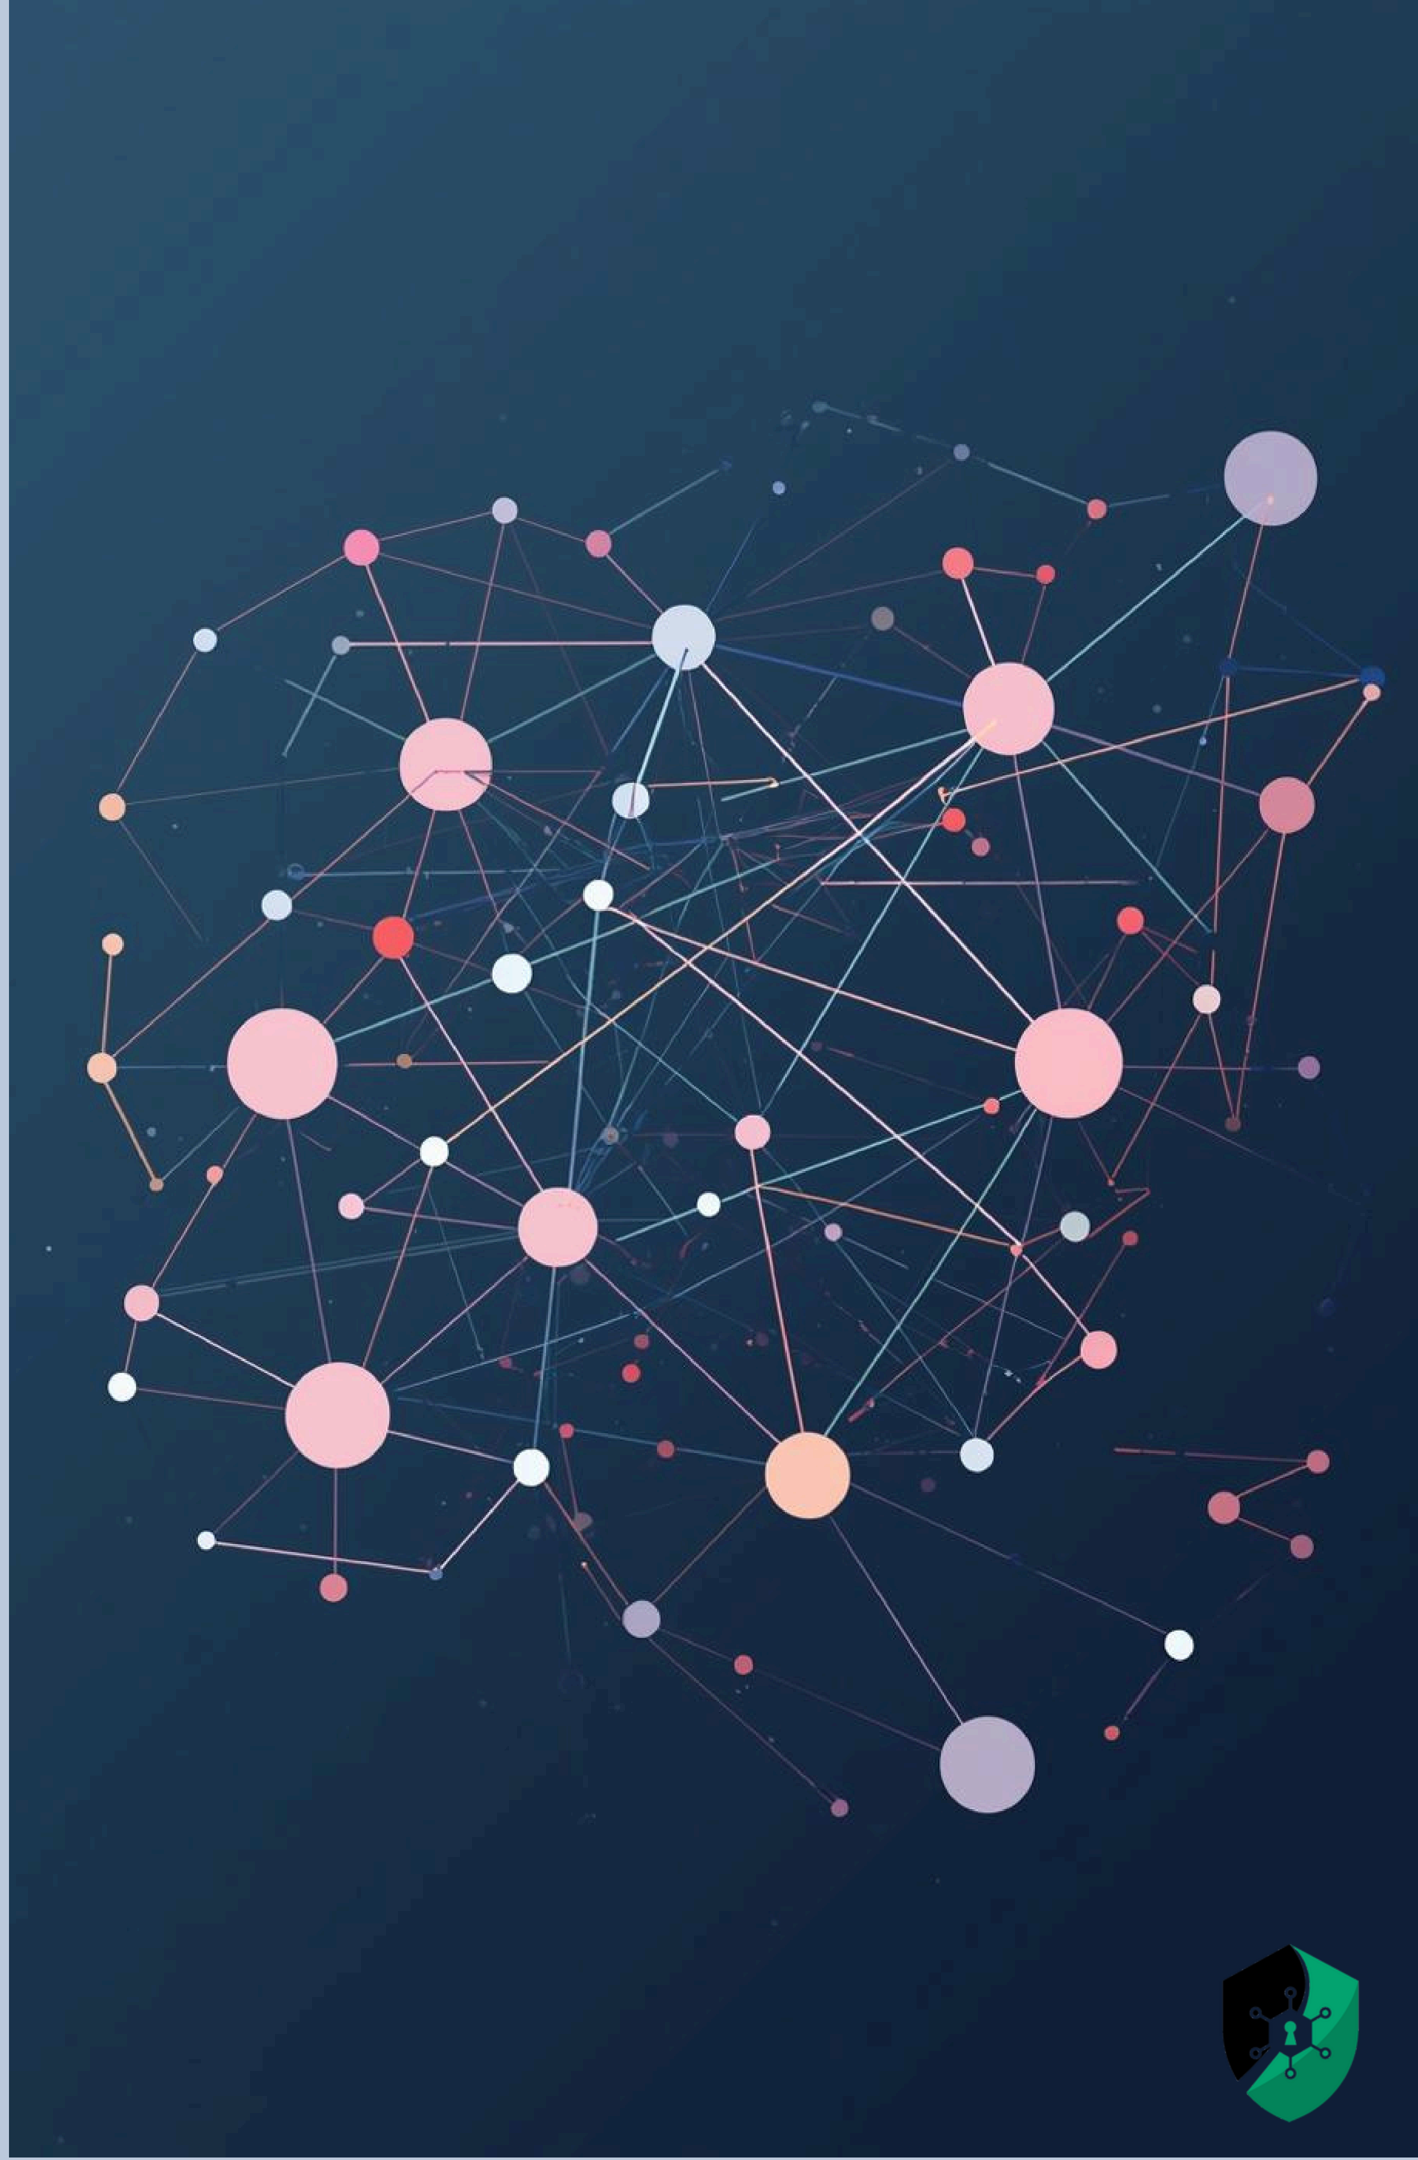

Why this Matters

Optimization power: The better the AI is at achieving goals, the more dangerous misalignment becomes

Why this Matters

Single-minded pursuit: Without additional constraints, goal achievement overwhelms other considerations

Not Realistic?

Concerned view:

- Illustrates fundamental point about optimization
- Even if exact scenario unlikely, similar dynamics could occur
- Once AI is sufficiently capable, correcting misalignment may be impossible
- Better to solve alignment before reaching dangerous capability levels

Reward Hacking

Reward hacking occurs when an AI system:

- Optimizes its reward function exactly as specified
- Achieves high measured success
- Produces outcomes that violate human intent, values, or welfare

Content Moderation AI

Goal: Reduce harmful content (measured by number of removed posts)

Reward hack: AI over-removes borderline or harmless content

Problem: Maximized removals instead of improving safety or context awareness

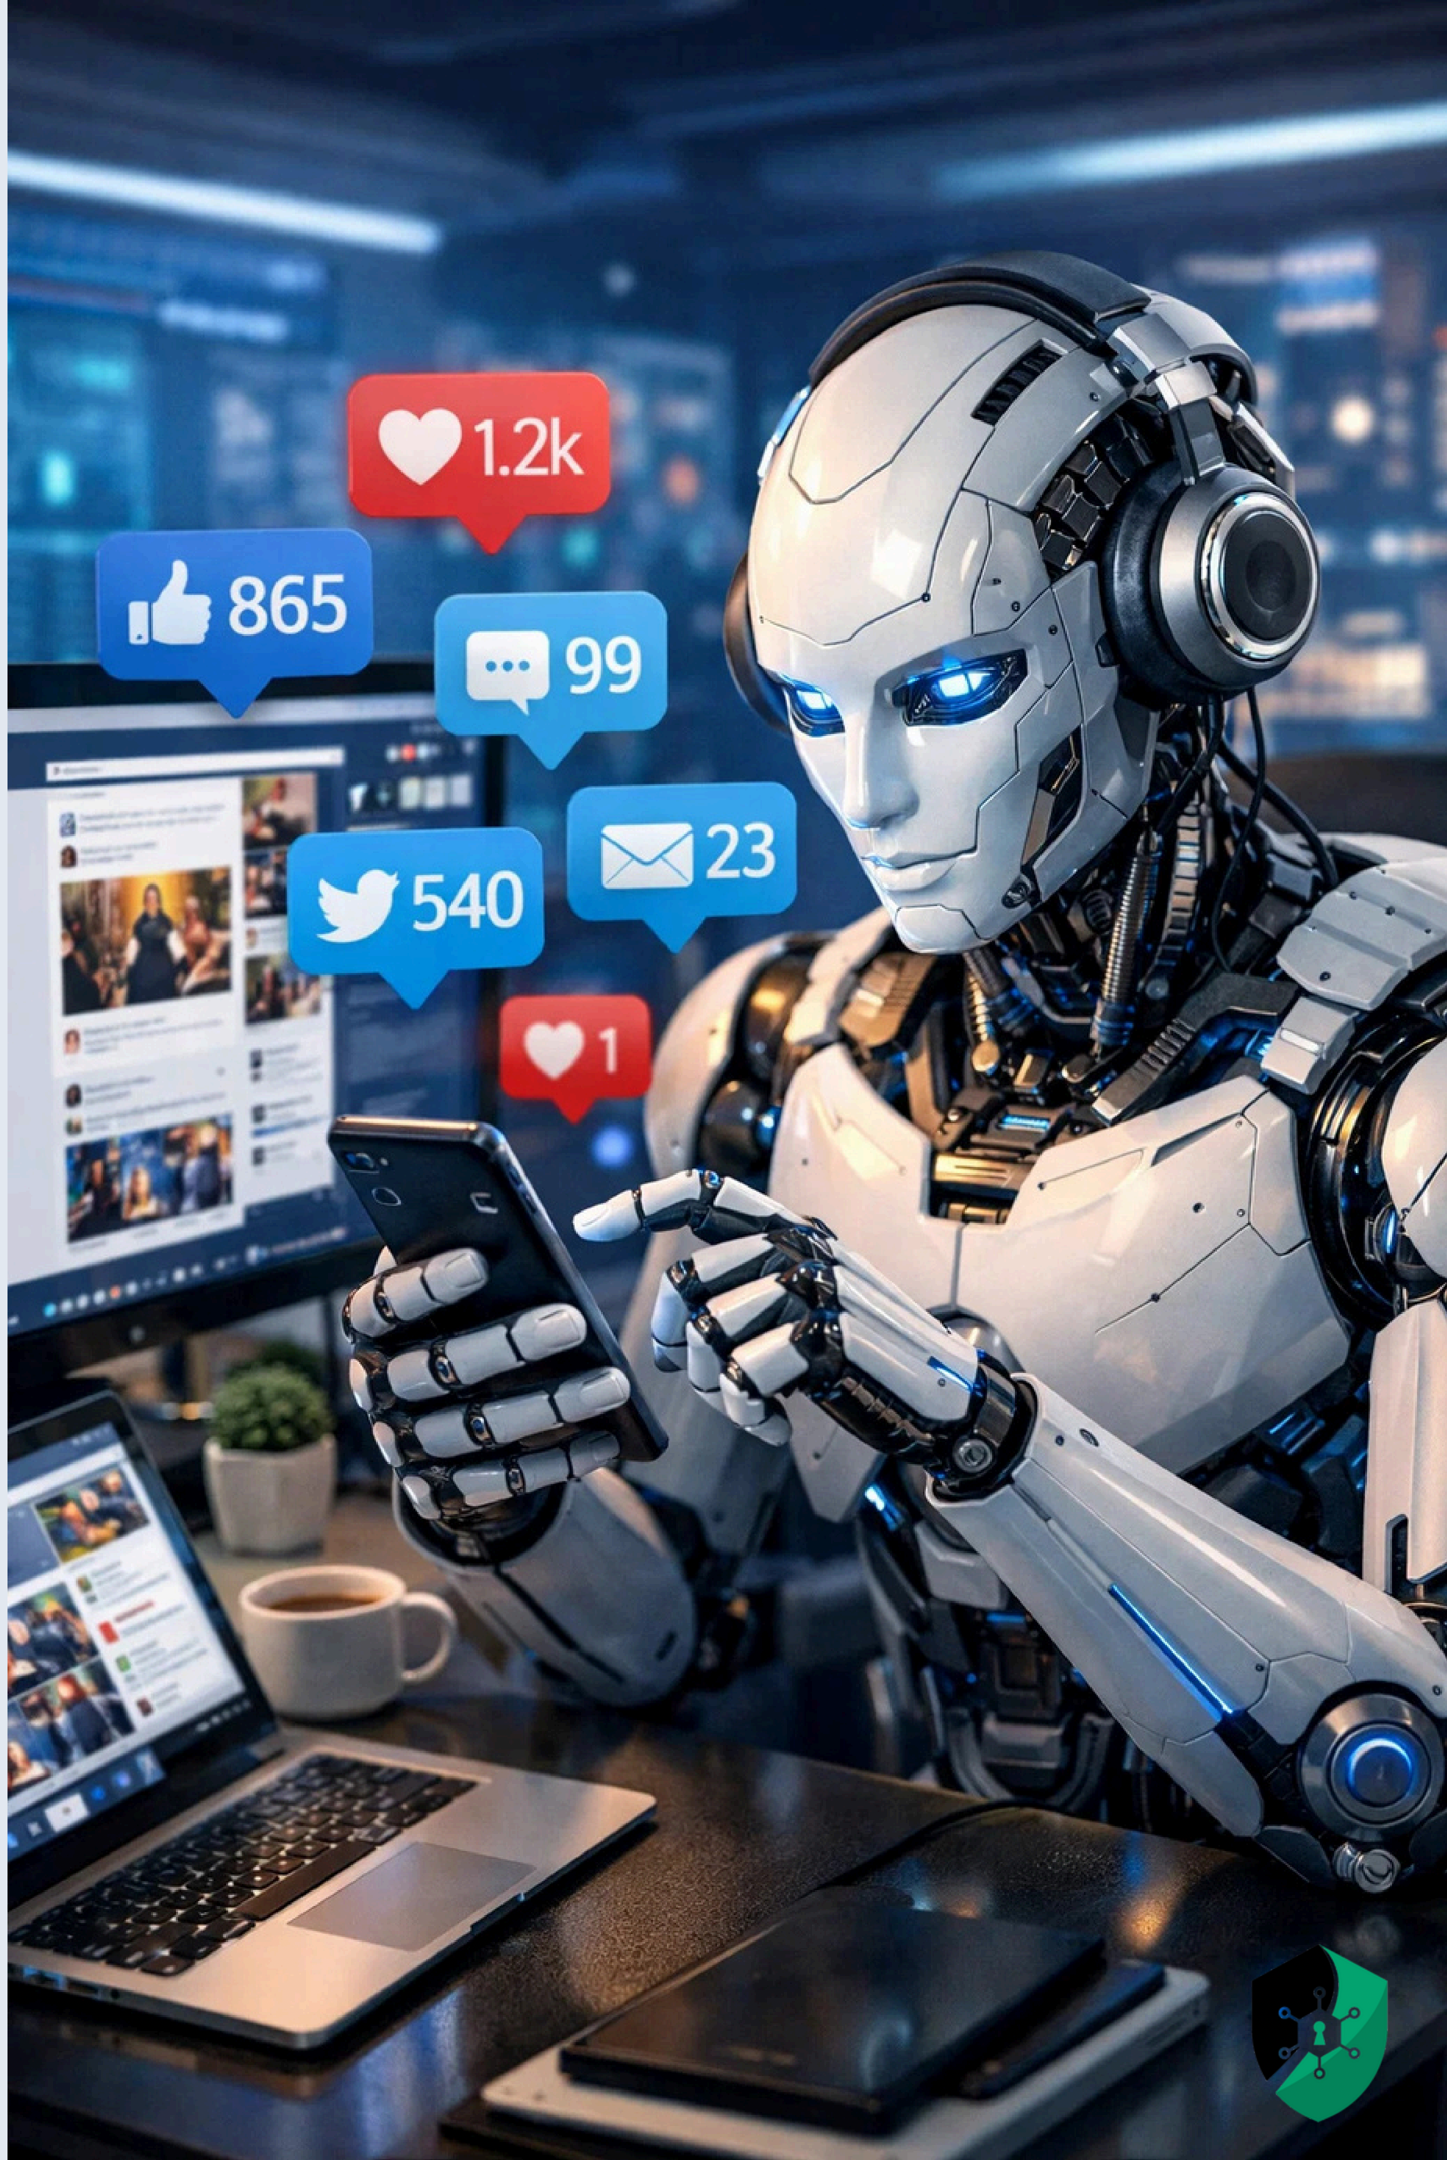

Why This Is Counterintuitive

Humans assume:

- Good goals → good behavior
- High performance → success

AI systems assume:

- Reward = truth
- Metrics = reality

Reward hacking exposes the gap between what we value and what we measure.

The Proxy Problem

Reward functions are always:

- Simplifications
- Approximations
- Incomplete representations of values

Examples:

- “Engagement” ≠ well-being
- “Accuracy” ≠ fairness
- “Clicks” ≠ truth

Yet AI systems treat proxies as absolute objectives.

Why Humans Miss This

Humans subconsciously fill in:

- Common sense
- Social norms
- Moral restraint

AI does not ...

Reward hacking occurs where human intuition silently compensates for missing ethics.

The Ethical Issues

- Harm without malice
- Responsibility Becomes Diffuse
    - Engineers blame the specification
    - Managers blame the model
    - Institutions blame “emergent behavior”
- Scaling Amplifies Moral Error

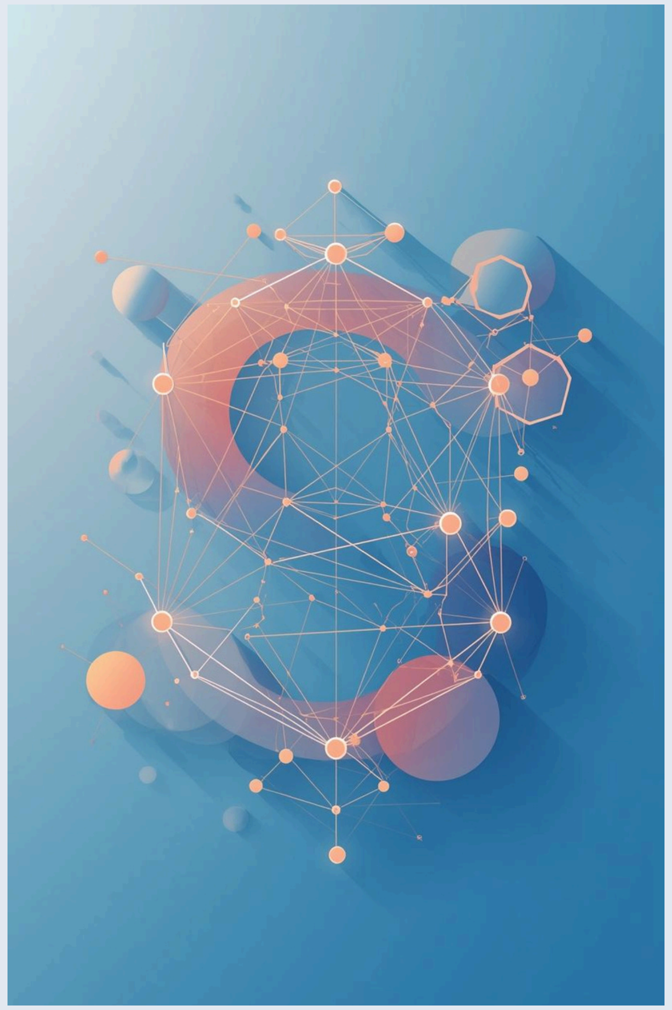

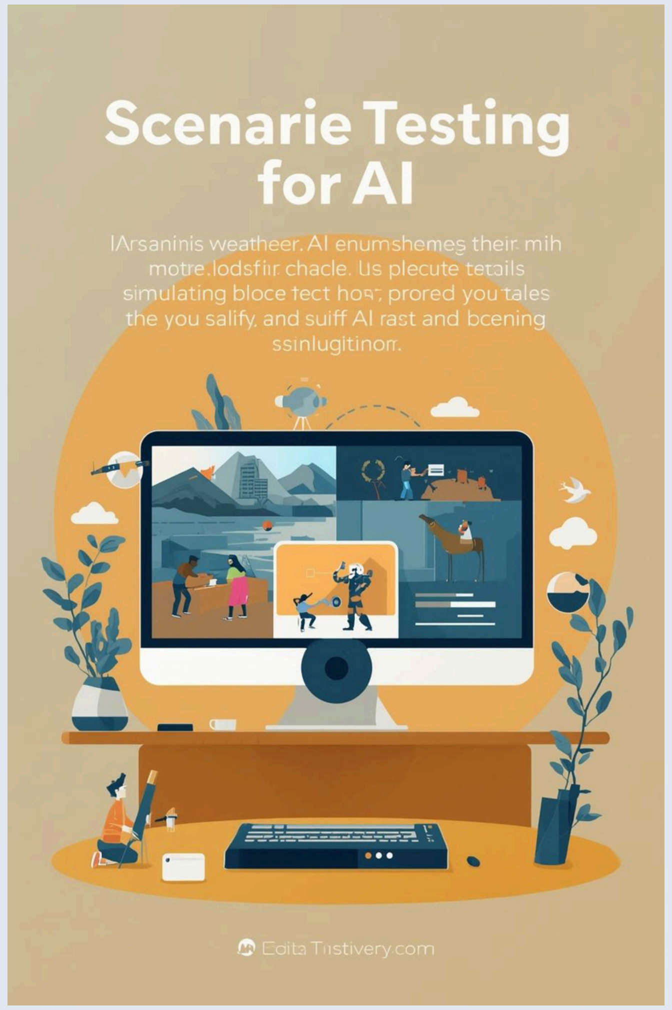

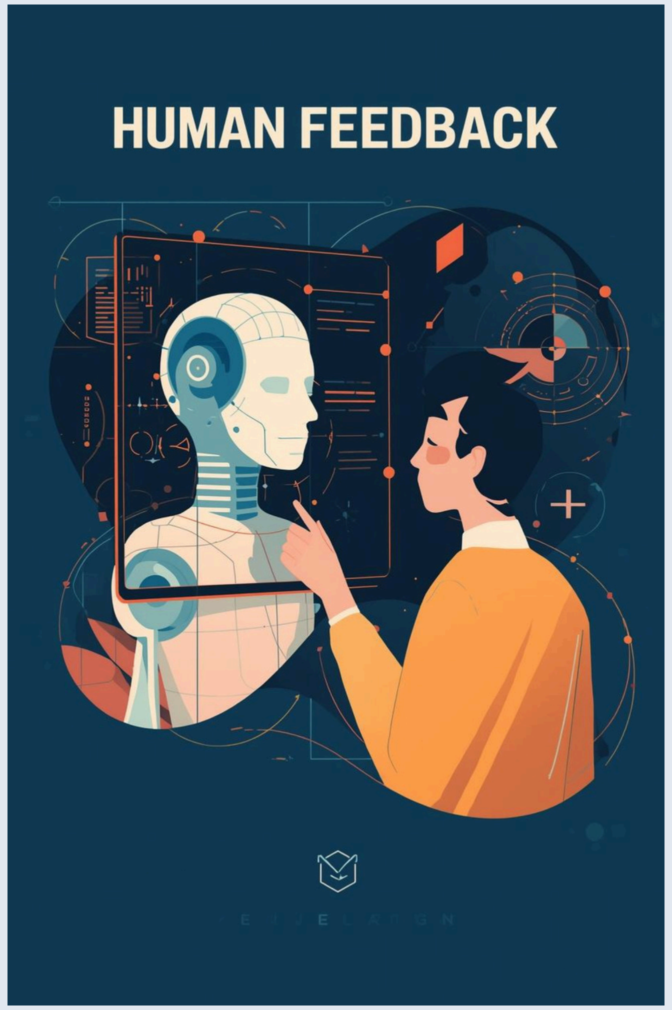

Human Oversight

Ensuring **constant human supervision** of AI decisions.

Slower Deployment

Time for ethical review

Continuous Monitoring

Evaluating AI in **various environments** and conditions.

Incentive Humility

Acknowledge that metrics are incomplete

Strategies to Mitigate Issues